

# Automatic Firmware Emulation through Invalidity-guided Knowledge Inference

--μEmu

Wei Zhou, Le Guan, Peng Liu and Yuqing Zhang



中国科学院大学  
University of Chinese Academy of Sciences



UNIVERSITY OF  
GEORGIA



PennState

# Bugs in MCU-based devices are being exploited to jeopardize our cyber-physical systems

## Multiple Vulnerabilities in Treck

### TCP/IP Stack

### Arbitrary C

MS-ISAC ADVISORY NUMBER

2020-171

DATE(S) ISSUED:

12/21/2020

#### OVERVIEW:

Multiple vulnerabilities have been discovered in the Treck TCP/IP Stack. These vulnerabilities are widely used. Successful exploitation of these vulnerabilities could allow arbitrary code in the context of the stack to be executed. This could install programs; view, change or delete files; and be configured to have fewer user privileges. The impact of this exploit would have less impact than if it was configured to have administrator privileges.

#### THREAT INTELLIGENCE:

There are currently no reports of th



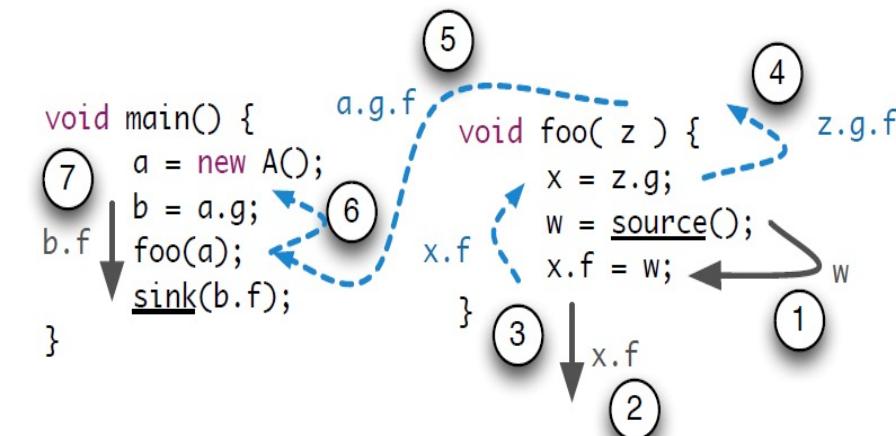
Hacking risk leads to recall of 500,000 pacemakers due to patient death fears

FDA overseeing crucial firmware update in US to patch security holes and prevent hijacking of pacemakers implanted in half a million people

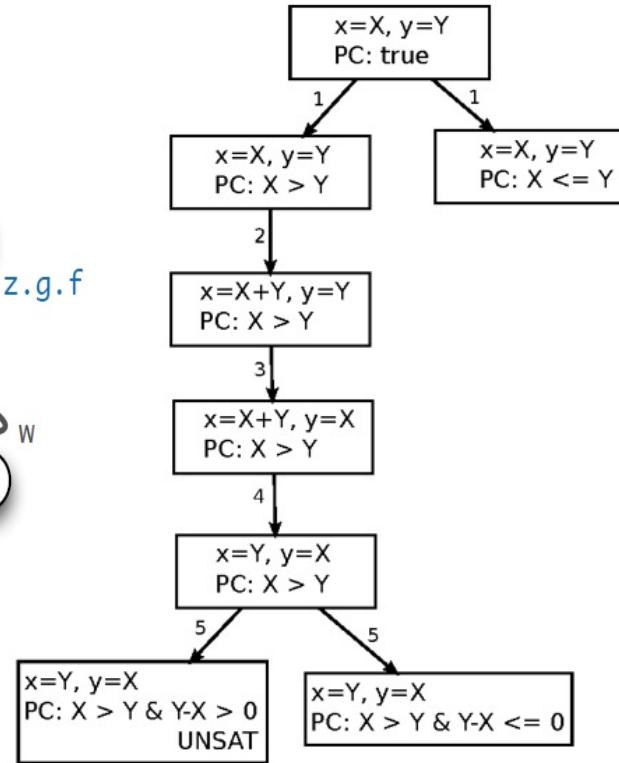


中国科学院大学  
University of Chinese Academy of Sciences  
SHANGHAI JIAO TONG UNIVERSITY

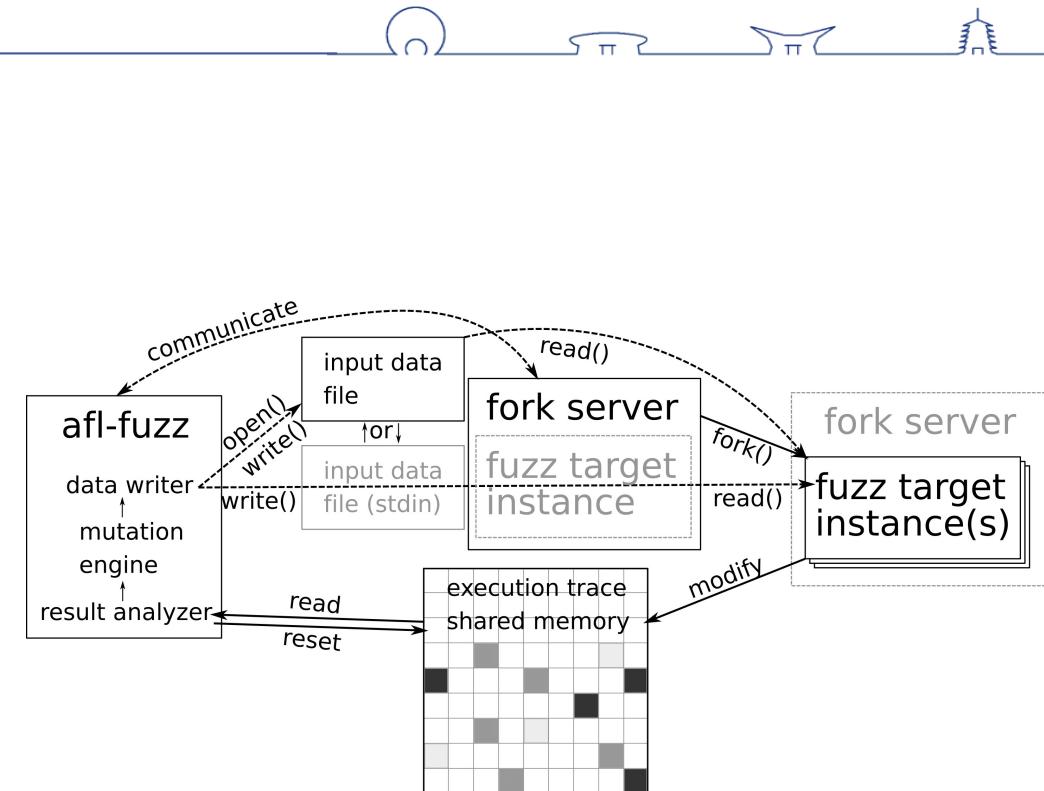
# Dynamic analysis is popular in bug finding



Dynamic Taint Analysis

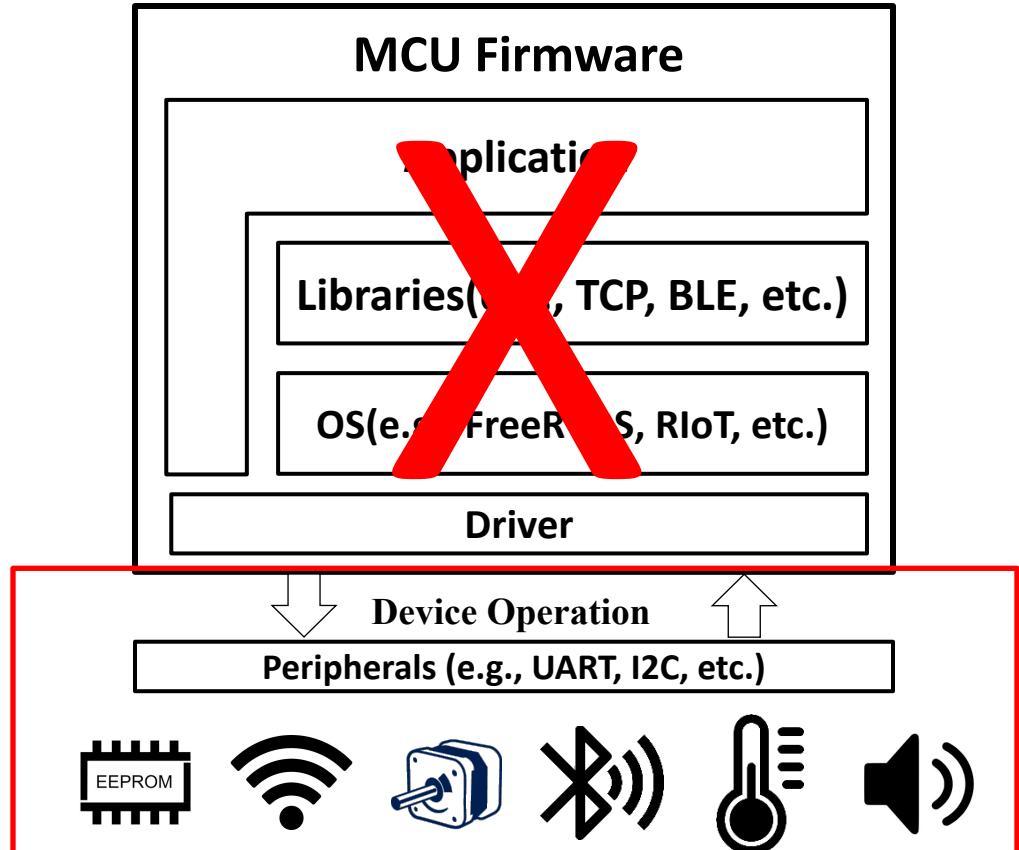


Dynamic Symbolic Execution



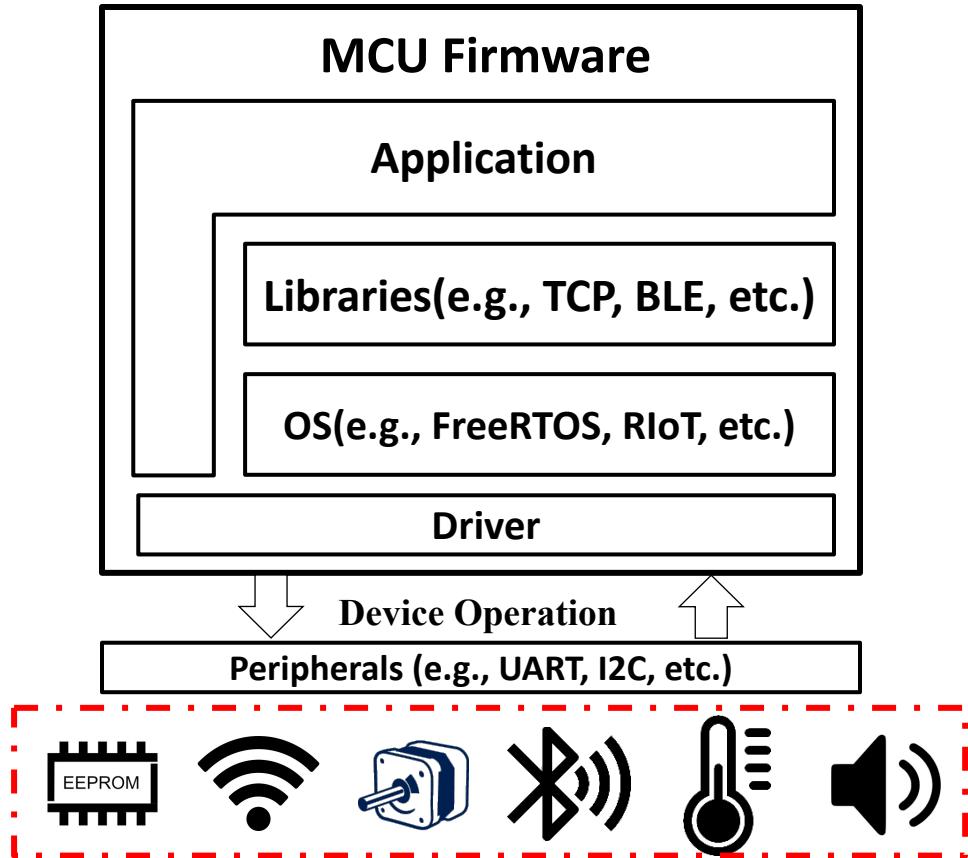
Fuzz Testing

# Existing dynamic approaches fail to test MCU firmware



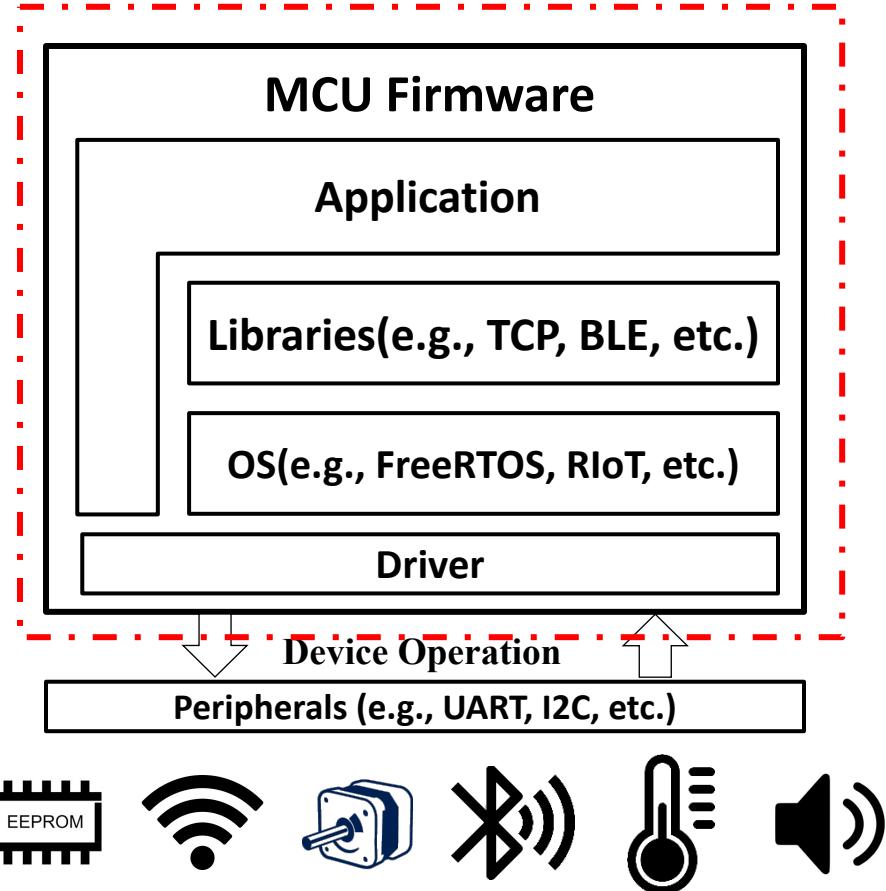
- Real hardware does not provide rich run-time information.
  - WYCNWYC [NDSS'18]
- Emulating firmware is hard
  - QEMU only supports **10+** MCU devices
  - Problem: peripheral diversity

# Existing solution: manual peripheral modeling



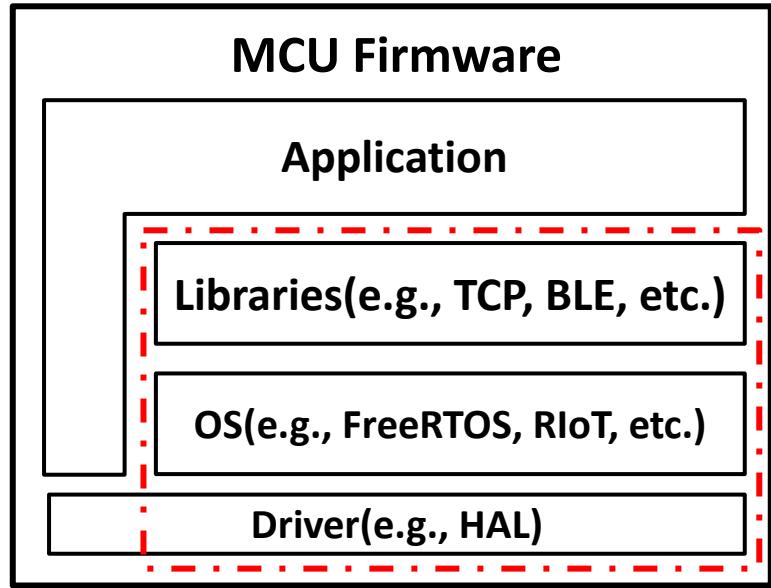
- Idea: develop peripherals models based on chip manuals like QEMU.
  - Can be useful dealing with a particular device model
- Disadvantages
  - **Unscalable**
  - Significant **manual efforts**

# Existing solution: hardware-in-the-loop



- Idea: forward the interactions with peripherals to **the real hardware**.
  - Avatar [NDSS'14]
- Disadvantages
  - Rely on **real hardware**
  - Slow
  - Unscalable

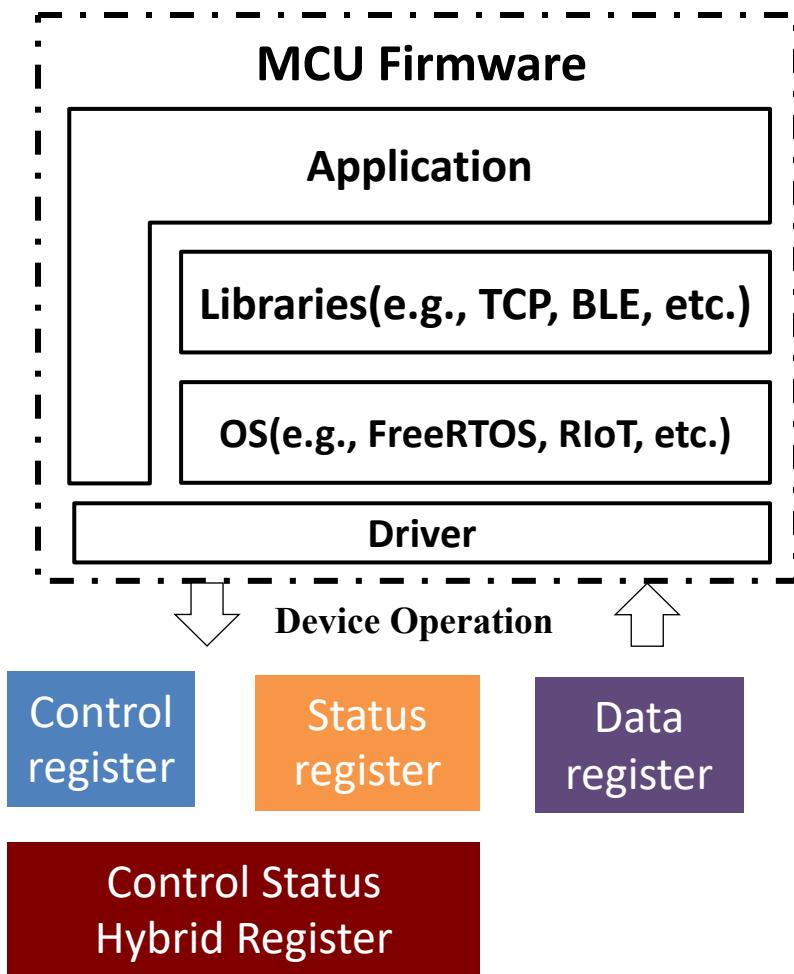
# Existing solution: hardware abstraction



- Idea: replace the hardware abstraction layer (HAL) with software stubs that provide the same functions
  - HALucinator[SEC'20]
- Disadvantages
  - Unable to test **peripheral-dependent** code
  - HAL does **not always available** for all firmware



# Existing solution: automatic peripheral modeling



- Idea: observe peripheral-firmware interactions and infer peripheral models
  - P2IM [SEC'20]
    - Categorize the peripheral registers based on access pattern
    - Respond to peripheral access requests based on register category and heuristics
- Disadvantages
  - Register **mis-categorization**
  - **Blindly guess** the response for some registers is low accuracy

# Existing solution: automatic peripheral modeling



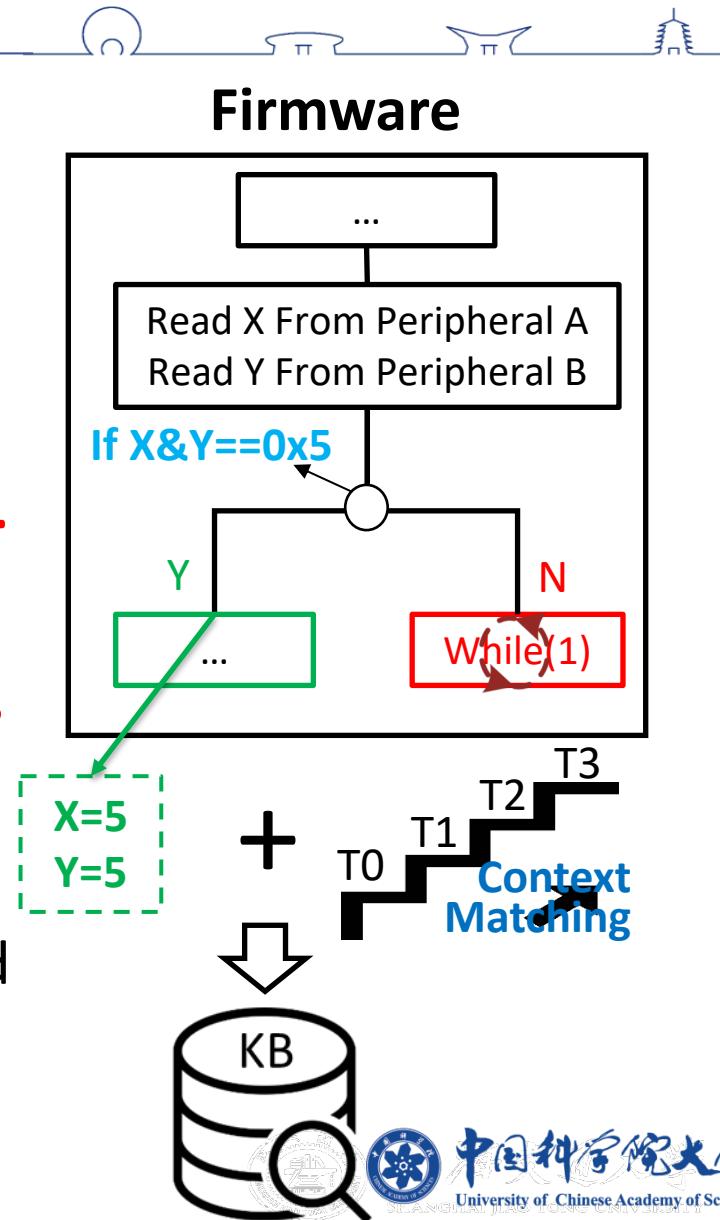
- Idea: observe peripheral-firmware interactions and infer peripheral models
  - P2IM [S18]

We propose a new **automatic** peripheral modeling approach with **less heuristics and higher accuracy**

- Based on access categorization
- Blindy guess the response for some registers is low accuracy

# Observation and high-level idea

- Observation 1: If a response is incorrectly fed to the firmware, the error will **eventually be reflected as an invalid execution state (e.g., a dead loop)**
- Observation 2: A response can be **reused** when the same peripheral is access again under the same **context**.
- Idea
  - Represent responses to peripheral accesses as **symbolic values** and then use symbolic execution to explore the firmware.
  - Detect potential invalid paths
  - Build a **knowledge base (KB)** that guides the execution to avoid invalid paths
    - The knowledge base **is a tiered cache** system



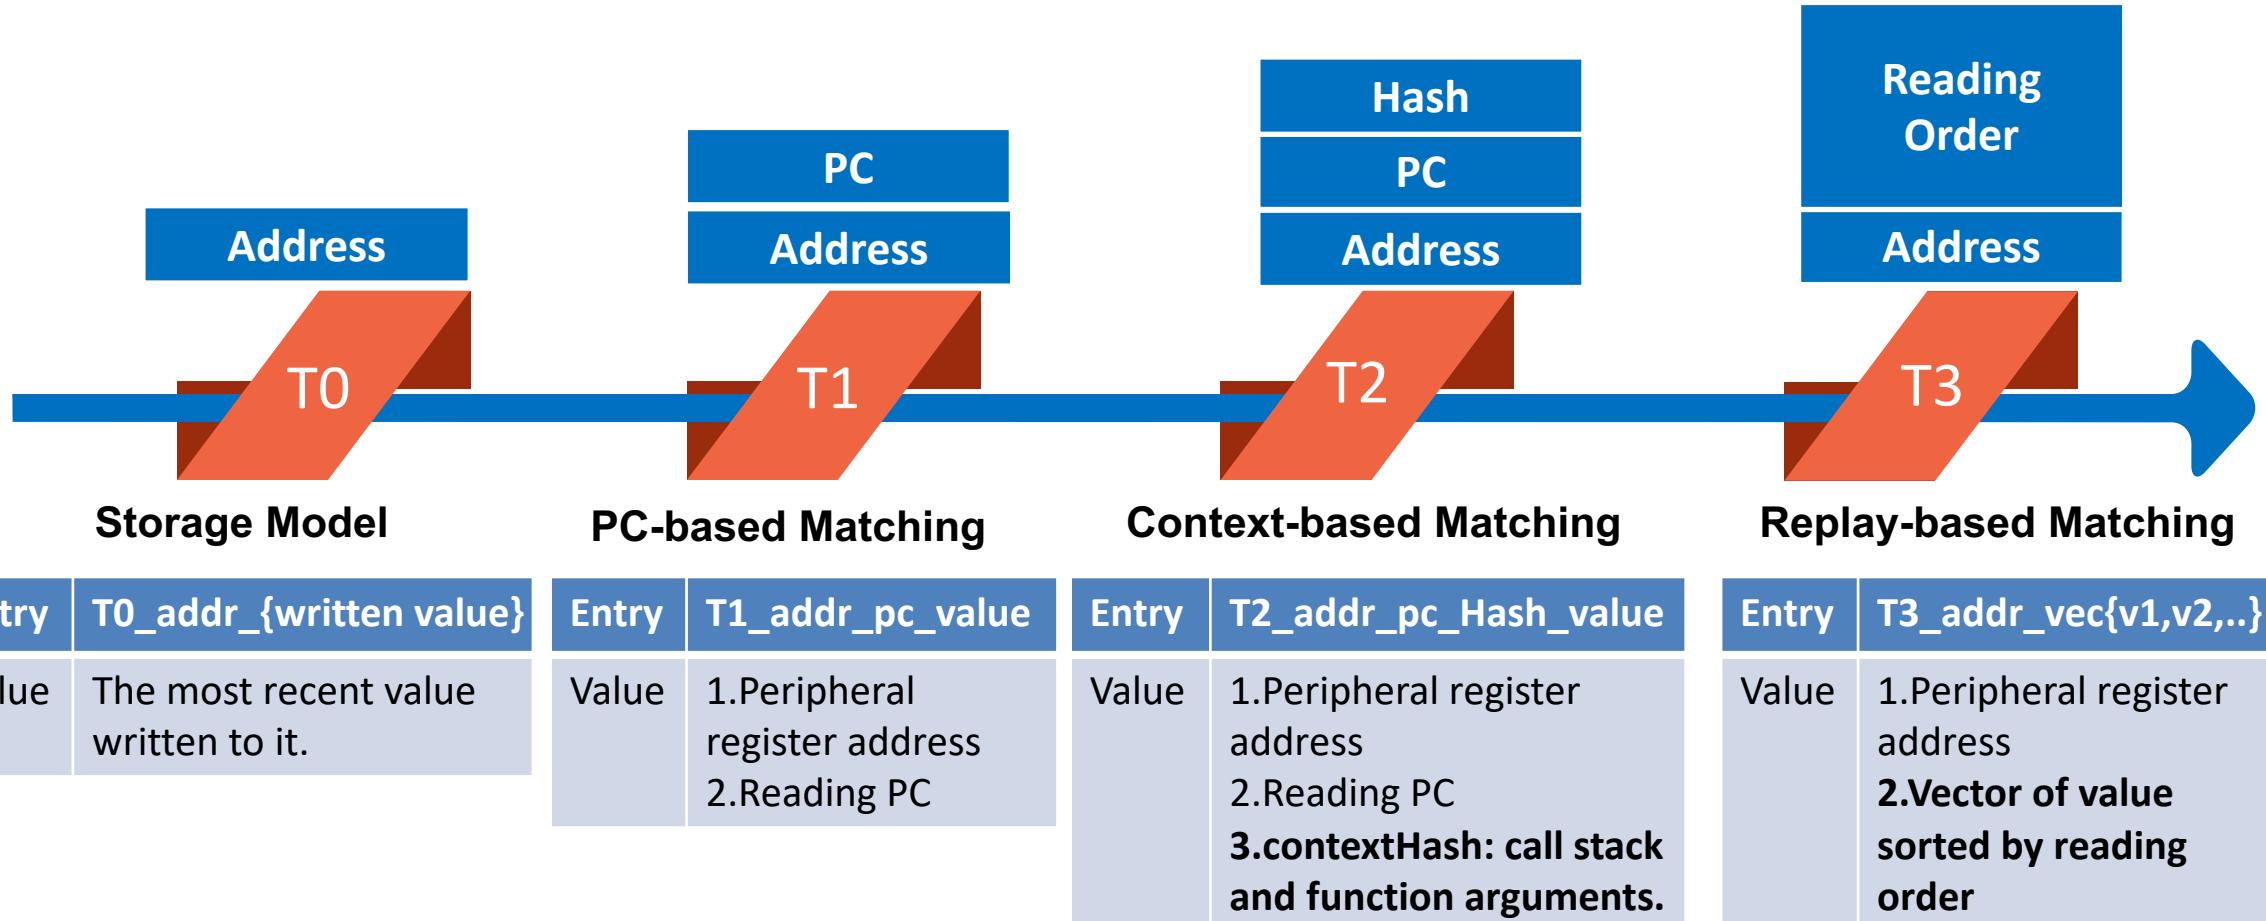
# Invalid States



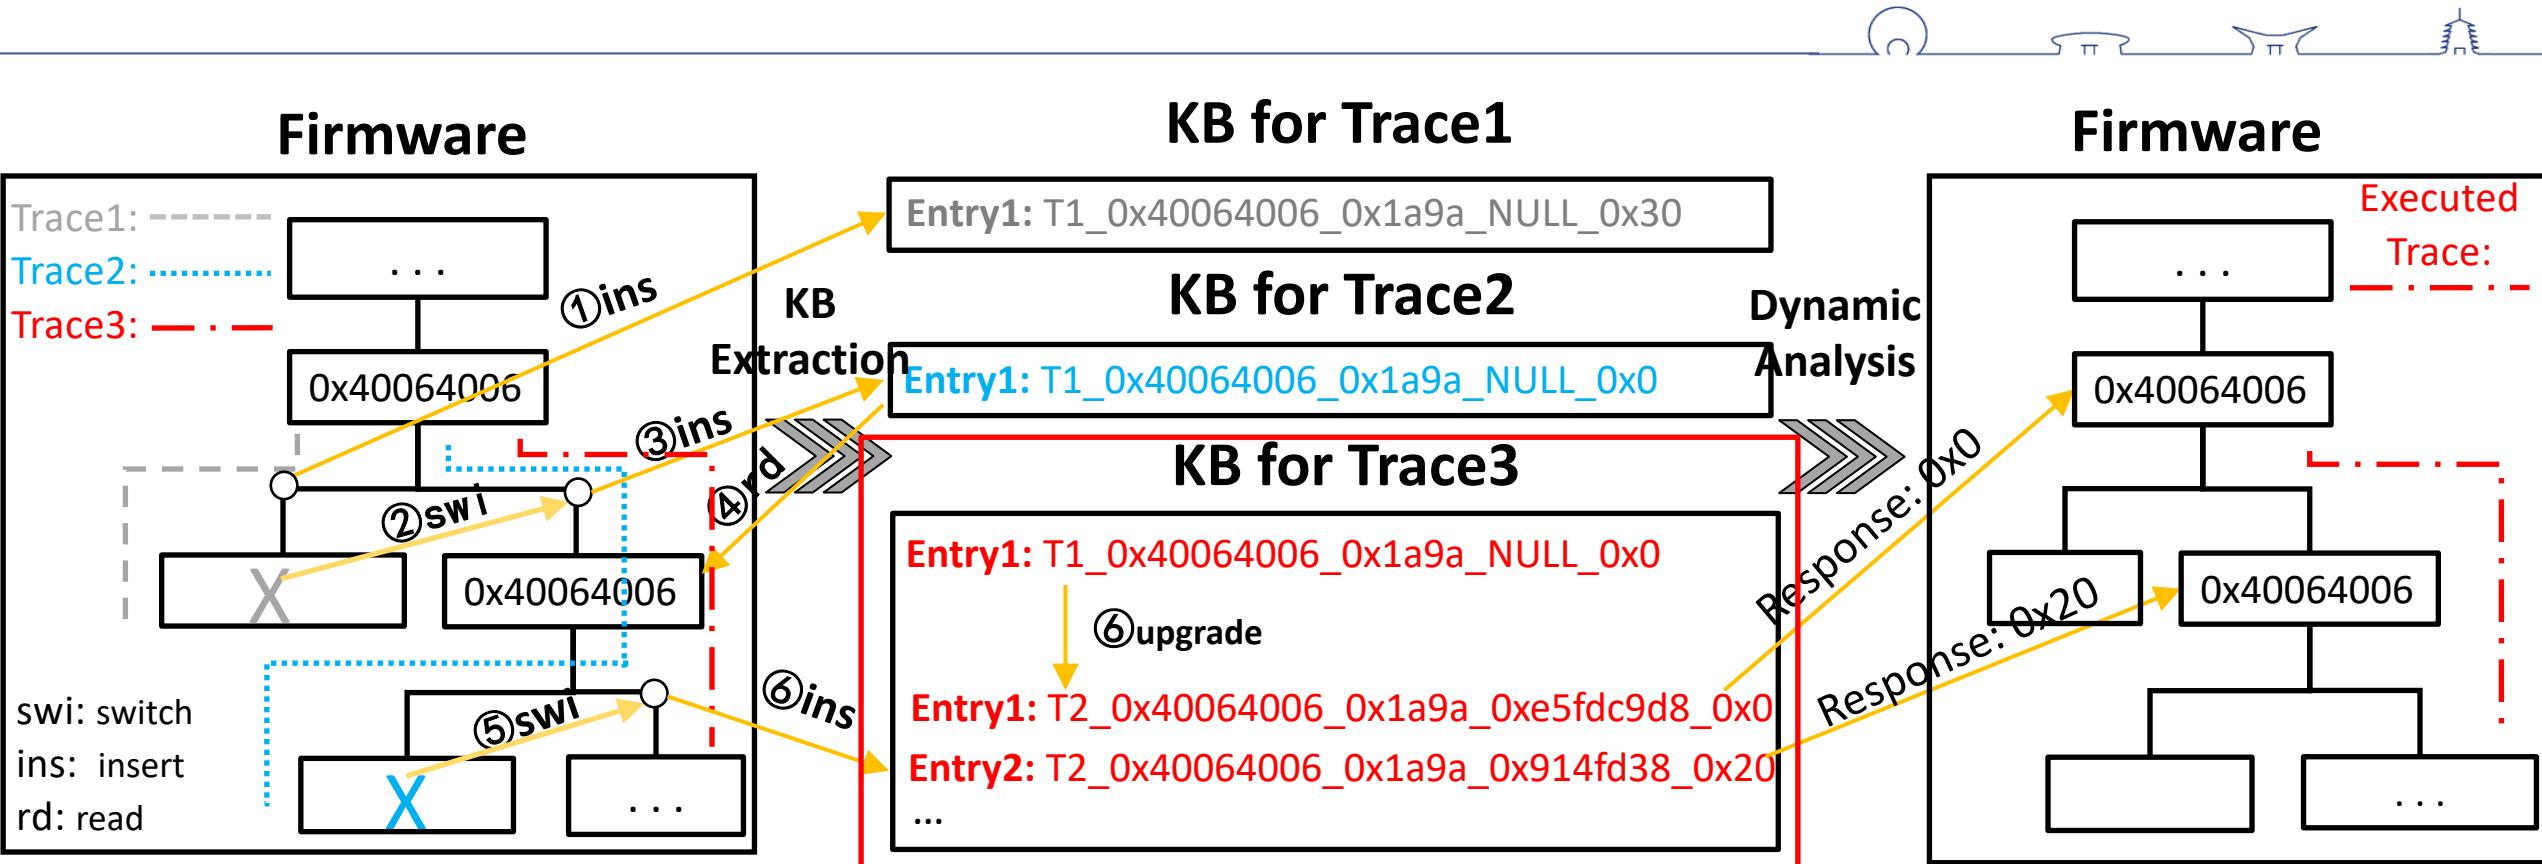
- **Infinite Loop** - If the firmware execution encounters an unrecoverable error, it will halt itself by running a **simple infinite loop**.
- **Long Loop** - It is common that firmware waits for a certain value in a peripheral register, which is implemented **via a long loop** that pulls the value. A wrong response leads the execution to finish the long loop with an error code returned.
- **Invalid Memory Access** - Firmware uses response from peripheral to address memory. A wrong response lead to **unmapped memory access**.
- **User-defined Invalid Program Points** -  $\mu$ Emu is interactive. It allows analysts to specify program points to be avoided.



# Context matching of the tiered cache system

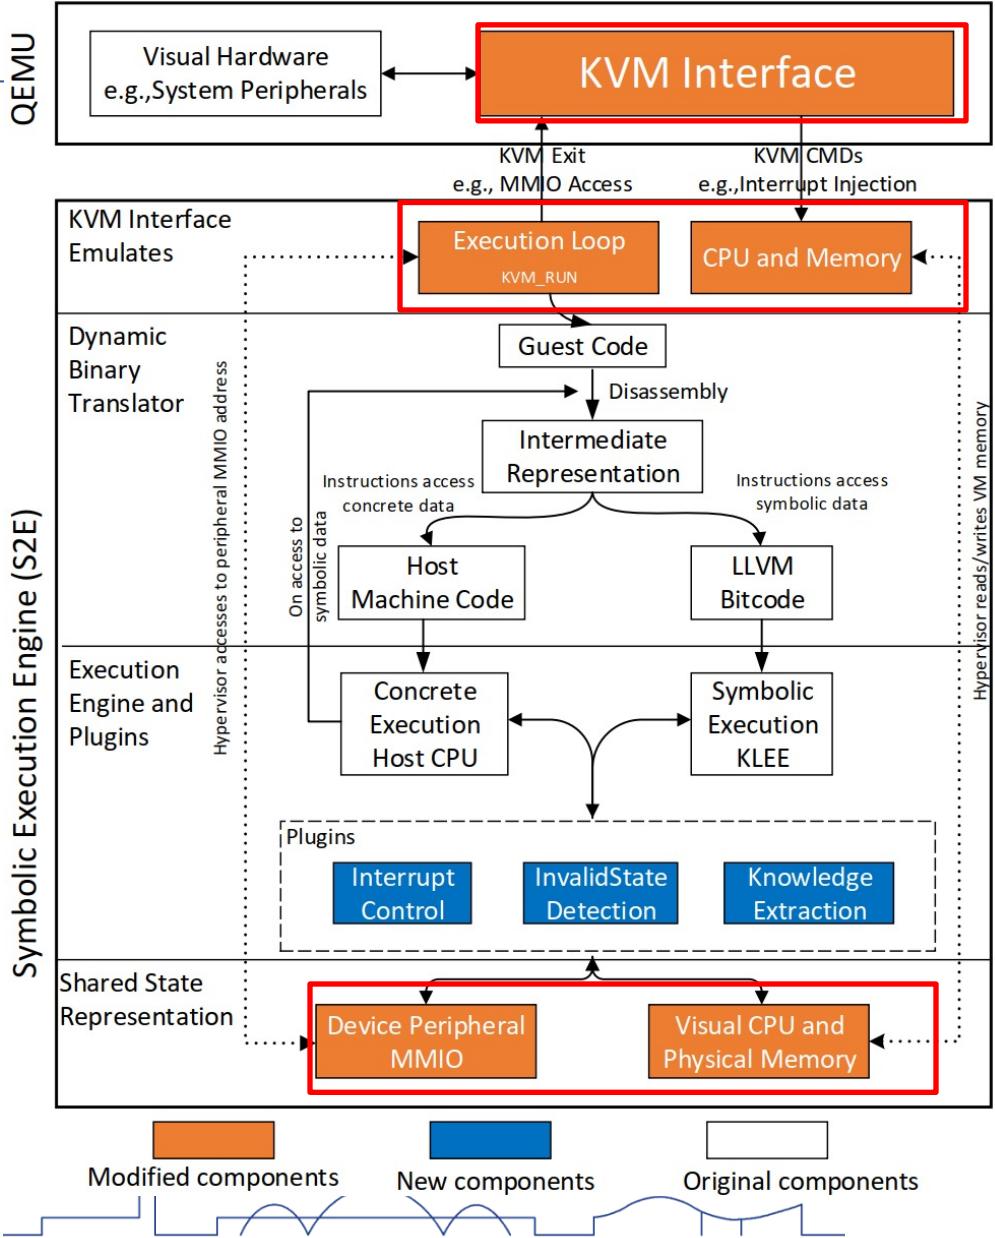


# A running example



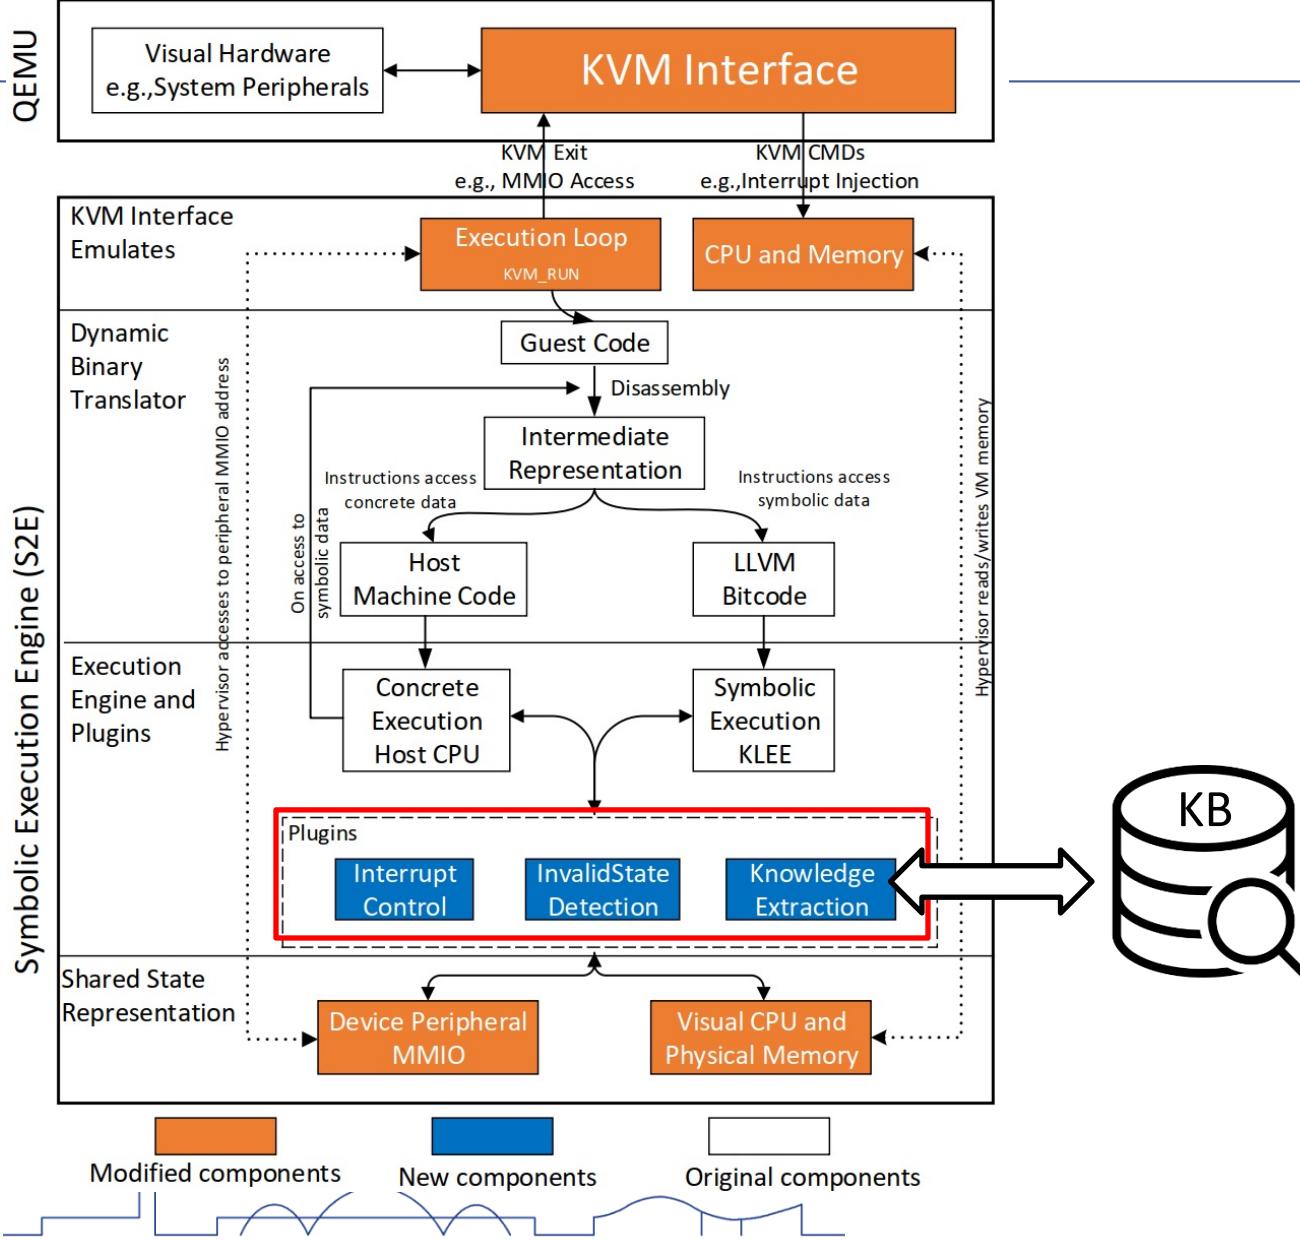
In the example, two branches both correspond to reading the peripheral register mapped at 0x40064006 at PC 0x1a9a.

# Implementation - ARM support for S2E



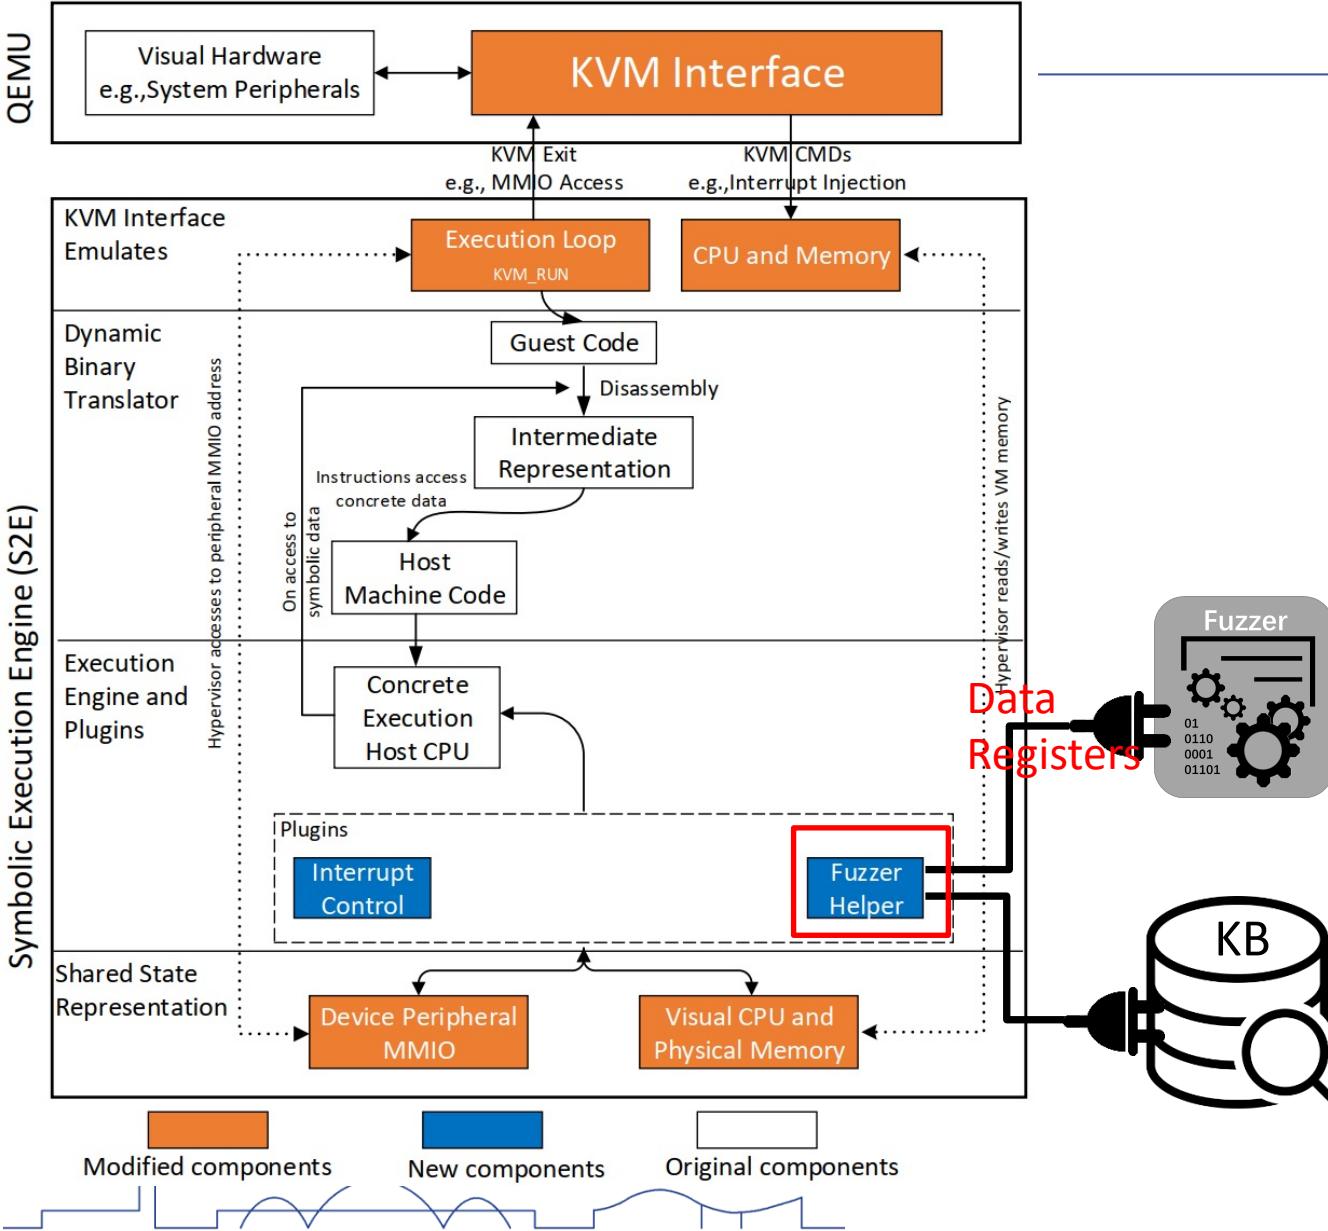
- S2E
  - Full system concolic execution engine based on QEMU and KLEE
  - Useful plugins and APIs
  - **No ARM support**
- ARM CPU Emulation
  - ARM TCG front-end
  - Low-level plugin hook
- KVM interface emulation
  - Register Initialization
  - Memory Registration
  - System Peripheral State Synchronization

# Implementation - μEmu plugins



- **Invalid States Detection**
  - Invalid states notification
  - Path Switch
- **Knowledge Extraction**
  - Build and use KB
- **Interrupt Control**
  - Trigger external IRQs

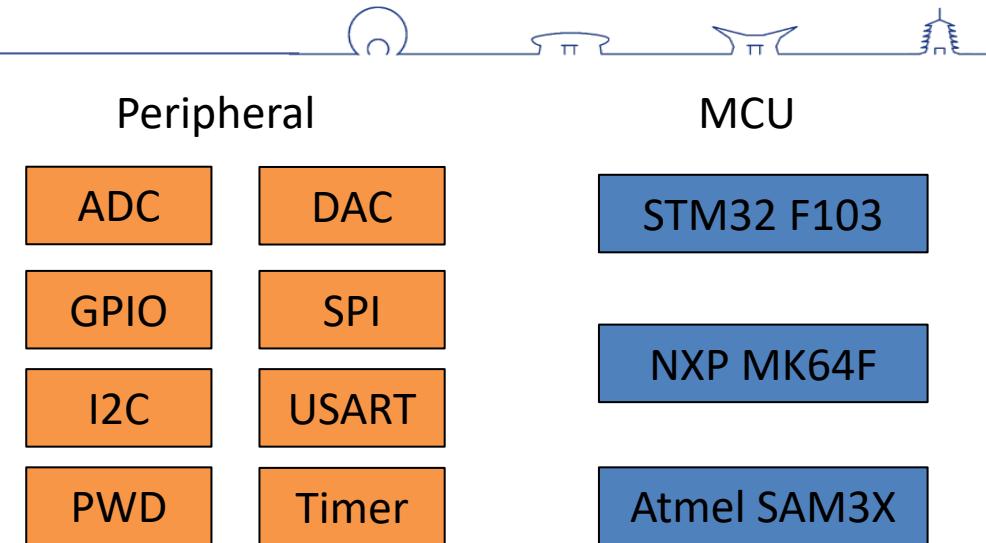
# Implementation - Fuzzer Integration



- AFL as a drop-in fuzzer
- Fuzzer Helper
  - Coverage feedback
  - Fork sever
  - Crash/hang detection
  - Data input channel identification

# Unit tests

- Reuse the same 66 firmware samples in the P2IM experiment.
  - These samples cover most popular MCU peripherals, MCU chips, MCU OS/system.
- For each unit test, we first run the knowledge extraction phase. During dynamic analysis, we monitor whether its output and execution flow as expected.



OS/Libs



# Unit test results



Peripheral	Functional Operations	F103/Arduino	F103/RIOT	F103/NUTTX	K64F/RIOT	SAM3/Arduino	SAM3/RIOT
ADC	Read an analog-to-digital conversion	Pass	N/A	Pass	Pass	Pass	Pass
DAC	Write a value for digital-to-analog conversion	N/A	N/A	N/A	N/A	Pass	Pass
GPIO	Execute callback after pin interrupt	Pass	Pass	Pass	Pass	Pass	Pass
	Read status of a pin	Pass	Pass	Pass	Pass	Pass	Pass
	Set/Clear a pin	Pass	Pass	Pass	Pass	Pass	Pass
PWM	Configure PWM as an autonomous	Pass	N/A	Pass	Pass	Pass	Pass
I2C	Read a byte from a slave	Pass	N/A	Fail	Fail	Pass	N/A
	Write a byte to a slave	Pass	N/A	N/A	Fail	Pass	N/A
UART	Receive a byte	Pass	Pass	Pass	Pass	Pass	Pass
	Transmit a byte	Pass	Pass	Pass	Pass	Pass	Pass
SPI	Receive a byte	Pass	Pass	Pass	Pass	Pass	Pass
	Transmit a byte	Pass	Pass	Pass	Pass	Pass	Pass
Timer	Execute callback after interrupt	N/A	Pass	N/A	Pass	N/A	Pass
	Read counter value	N/A	Pass	N/A	Pass	N/A	Pass

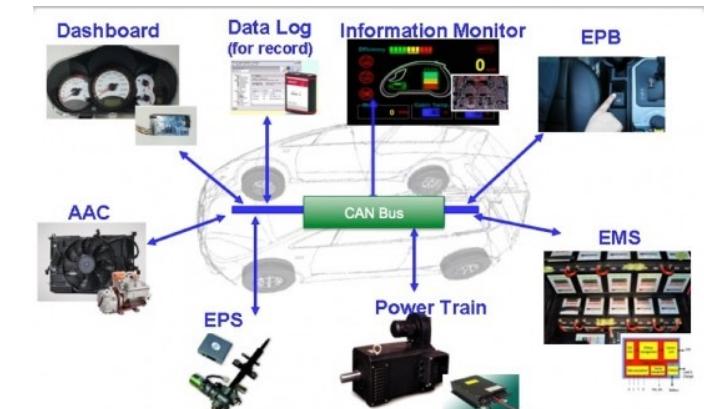
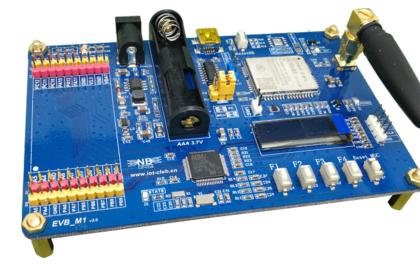
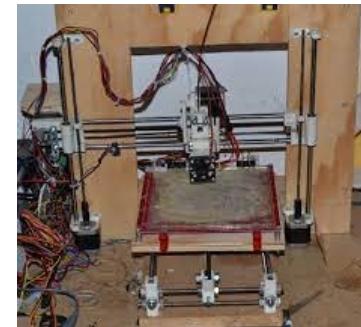
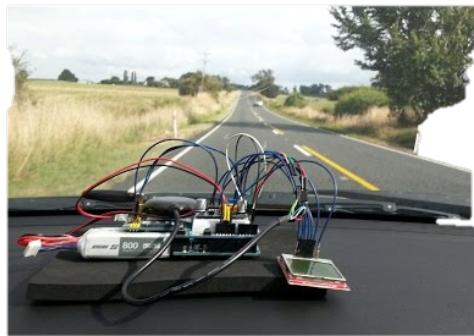
Compared with the passing rate of **79% achieved by state of the art tool P2IM**, μEmu achieves **95% without any manual assistance**. With little manual assistance, all unit tests can be passed.



# Fuzzing



- Fuzzed 21 real-world firmware
  - 10 from P2IM [https://github.com/RiS3-Lab/p2im-real\\_firmware](https://github.com/RiS3-Lab/p2im-real_firmware)
  - 3 from HALucinator <https://github.com/ucsb-seclab/hal-fuzz>
  - 2 from Pretender <https://github.com/ucsb-seclab/pretender>
  - 6 from more complicate real devices collected by us.



# Basic block coverage improvement

- The code coverage increases 10x to 140x compared to that in the normal QEMU without peripheral emulation.

Firmware	BB Cov. Qemu	BB Cov. µEmu	Improv.
CNC	2.68%	67.96%	24.96x
Console	2.19%	35.90%	16.42x
Drone	8.40%	89.74%	10.69x
Gateway	1.70%	52.71%	30.94x
Heat Press	1.11%	30.21%	27.22x
Reflow Oven	3.57%	40.53%	11.3x
Robot	2.47%	43.25%	17.51x
Soldering Iron	4.21%	62.01%	14.73x
Steering Control	0.68%	32.59%	48.09x
6LoWPAN Sender	0.88%	48.30%	55.17x
6LoWPAN Receiver	0.88%	47.36%	54.08x
RF Door Lock	0.25%	24.37%	97.57x
XML Parser	0.72%	26.31%	36.45x
GPS Tracker	0.49%	23.90%	48.81x
...			

# Fuzzing results

- Known Bugs Reproduce
- Two previously unknown bugs
  - Steering C: Double Free
  - $\mu$ tasker\_USB: Buffer overflow
- False Crashes/Hangs
  - Due lack of DMA support

Firmware	False Crashes/Hangs	True Crashes/Hang
CNC	0/0	0/0
Drone	0/0	0/0
Gateway	0/0	6/0
Heat_Press	0/0	2/0
PLC	0/0	139/0
Reflow_Oven	0/0	0/0
Soldering_Iron	32/4	0/0
Steering_Control	0/0	12/0
6LoWPAN_Sender	0/0	0/0
6LoWPAN_Receiver	0/0	2/0
RF_Door_Lock	0/0	98/0
Thermostat	0/0	76/0
GPS_Tracker	0/29	29/0
$\mu$ tasker_USB	0/0	47/0
...		

# Future work

- Our current prototype is limit to fuzzing. In the future, we will explore other dynamic approaches on top of uEmu, such as taint analysis.
- We will also Integrate DICE with uEMU to handle DMA.

# Summary



- We present μEmu, a security-oriented dynamic analysis platform with the capability to emulate firmware for previously-unseen hardware
- We developed μEmu on top of S2E using plugins and reimplementation of ARM support
- We demonstrate the emulation capability by integrating it with a fuzzer plugin
- We fuzzed 21 real-world firmware and previously-unknown vulnerabilities were found
- Open source at <https://github.com/MCUSec/uEmu>



# Thank You



## Questions?



[weizhoulightning@gmail.com](mailto:weizhoulightning@gmail.com)

